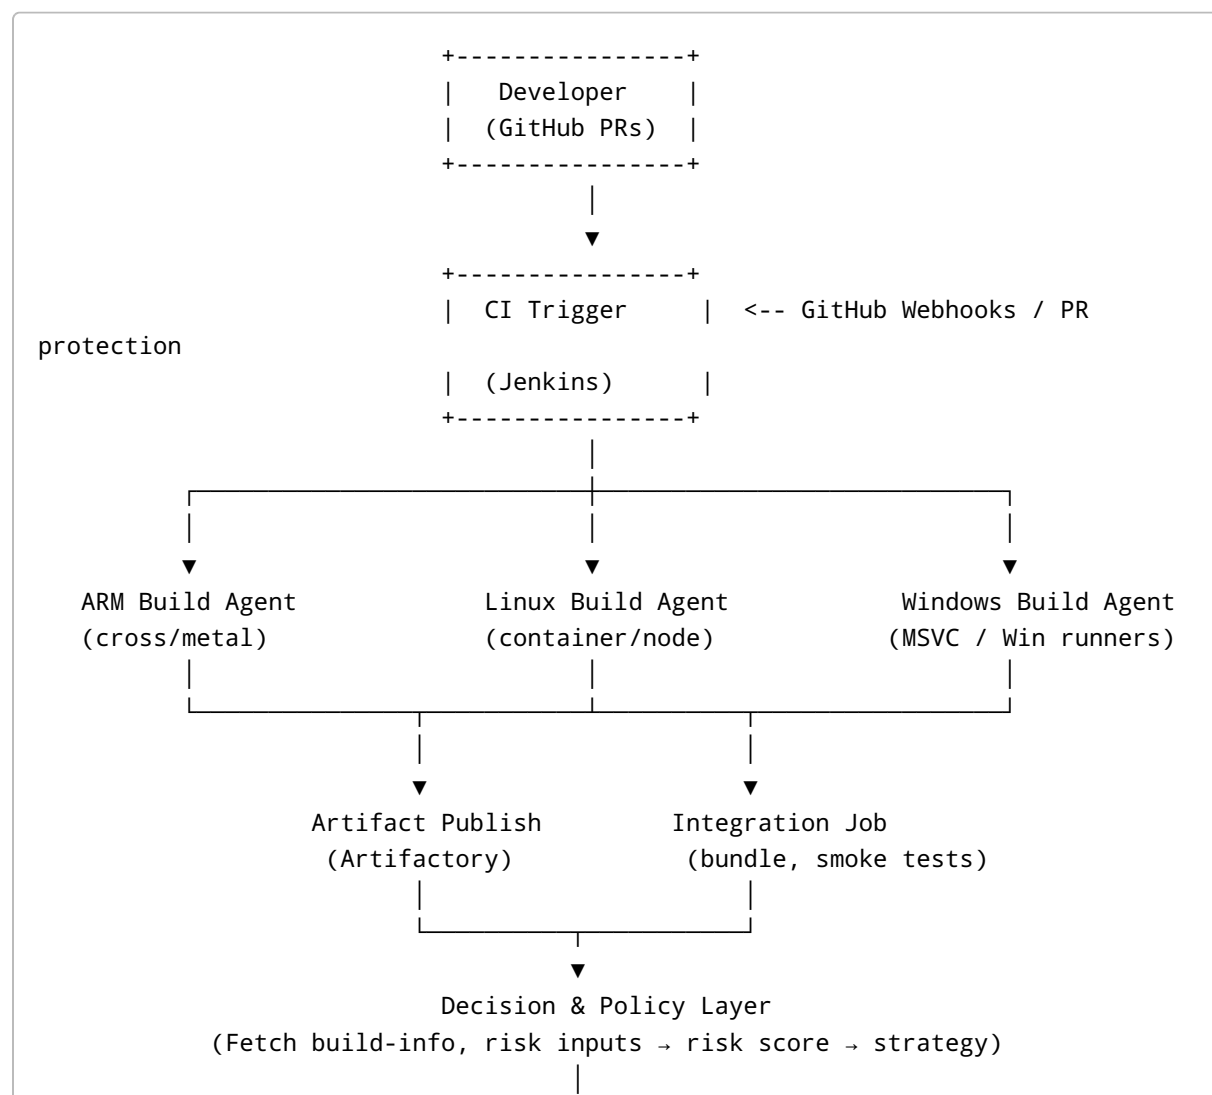


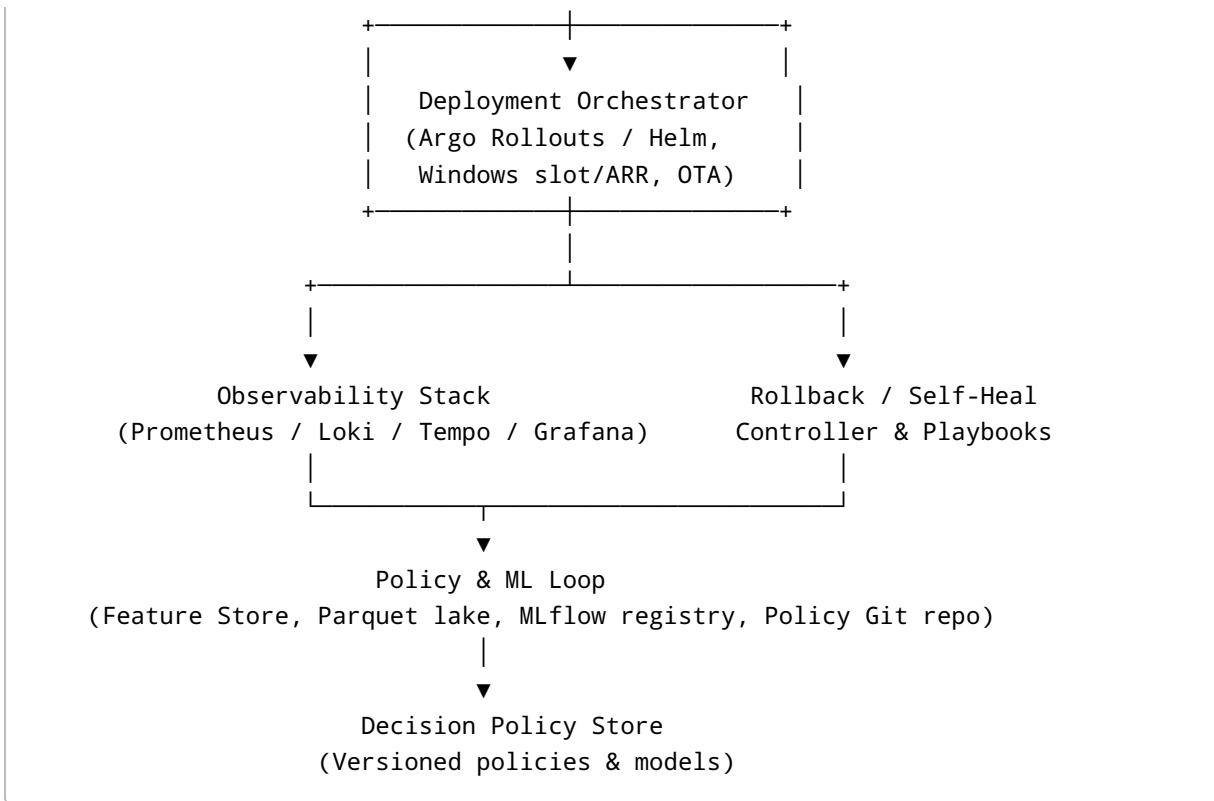
End-to-End Architecture: AI-Driven Adaptive CI/CD (Jenkins + Multi-OS)

This document describes a complete, end-to-end architecture for the pipeline we built out: from developer push → multi-OS builds (ARM/Linux/Windows) → artifact publishing → AI-driven risk scoring → strategy selection (Rolling/Canary/Blue-Green) → progressive deployment + real-time monitoring → auto-rollback/self-heal → continuous learning and policy improvement.

It is written as a clear architecture diagram (ASCII + sections), component responsibilities, data flows, control flows, and operational notes so you can implement it in your organization.

1) High-level ASCII Architecture





2) Key Components & Responsibilities

- **Developer / GitHub:** Authors of C++ repos (drivers, firmware, user-space). Triggers CI via PR/push.
- **Jenkins (CI):** Multibranch pipeline that:
 - checks out code, runs linters & static analysis
 - runs per-OS build stages in parallel (ARM/Linux/Windows)
 - runs unit and emulator/hardware-in-the-loop smoke tests
 - produces `build-info.json` and `risk-inputs.json`
 - publishes artifacts to Artifactory under `candidates/` with metadata
- **Artifactory:** Stores artifacts + attached metadata (build-info, SBOM, test reports, decision audit links). Holds `latest` and `candidate` namespaces.
- **Integration Job:** Bundles multi-repo artifacts (driver+firmware) and runs cross-repo integration smoke tests in isolated Green-like environment.
- **Decision & Policy Layer:**
 - Fetches `risk-inputs.json`

- Loads policy bundle (weights or model) from Policy Git or Policy Service
- Computes `risk_score` (heuristic or ML), produces `risk-decision.json`
- Persists decision and policy version to Artifactory

- **Deployment Orchestrator:**

- K8s: Argo Rollouts / Flagger + Helm manifests
- Windows: IIS slot/pool swap + ARR weight control
- ARM: OTA manager (Mender/Balena/custom) and device cohort manager
- Consumes `risk-decision.json` to select Rolling/Canary/Blue-Green flow

- **Observability Stack:**

- OpenTelemetry client libs in apps
- Prometheus (metrics) + recording rules
- Loki (logs) + Tempo/Jaeger (traces)
- Grafana dashboards for Release Compare + SLOs + Canary progress

- **Analyzer (AIOps):**

- Runs EWMA / Z-score / STL / Isolation Forest analyses on candidate vs baseline
- Produces `verdict` (promote|slow|rollback) and `evidence` (metric deltas, plots)

- **Rollback & Self-Heal Controller:**

- Executes idempotent playbooks for K8s/Windows/ARM
- Locks execution, records action, triggers incident creation if needed

- **Policy & ML Loop:**

- Feature store (Feast or parquet lake) collects build/test/deploy outcomes
- MLflow tracks models; models predict failure probability or recommend tests
- Policies (YAML) stored in Git with review process

3) Data Flow (detailed)

1. Developer pushes commit → Jenkins pipeline triggered.
2. Jenkins builds per-OS, runs tests, collects SBOM & coverage → creates `build-info.json` & `risk-inputs.json`.
3. Jenkins publishes artifacts to Artifactory under `candidates/` and writes metadata.
4. Integration job assembles bundle (if needed) and publishes integration manifest.

5. Decision layer fetches metadata + current policy bundle → computes `risk_score` and `deploy_strategy`.
 6. Decision stored in Artifactory and posted to Slack; pipeline branches to selected deployment path.
 7. Orchestrator deploys candidate to Green / runs canary weights / executes rolling batches.
 8. Observability collects metrics/logs/traces; Analyzer compares candidate vs baseline and emits `verdict`.
 9. If `verdict==promote` → orchestrator increases rollout; if `slow` → extend watch; if `rollback` → rollback controller invoked.
 10. All decisions & evidence saved; feature store updated for ML training.
 11. ML/training jobs periodically run on collected data to update policies/models; policies versioned in Git and rolled out in shadow first.
-

4) Control Flow (decision & safety)

- **Policy enforcement levels:**
 - Shadow: decisions logged but not enforced (collect counterfactuals)
 - Advisory: decisions posted to Slack/PR but require manual approval
 - Enforced: automation acts on decision (deploy path controlled)
 - **Manual override:** pipeline supports manual approvals for high-risk changes (DB change, kernel driver, critical CVE).
 - **Cooldown:** after an automated rollback, that candidate is quarantined for a configurable cool-down window.
-

5) Per-OS Specific Patterns & Safe Defaults

- **ARM (Firmware):**
 - Always start Green on small cohort of test devices. Prefer canary cohorts; Blue/Green for risky releases.
 - Firmware images immutable and signed; previous image retained for instant rollback.
 - Extra telemetry for device health (boot time, crash beacon).
 - **Linux (K8s / Containers):**
 - Use Argo Rollouts / Flagger for progressive delivery. Argo AnalysisTemplate uses Prometheus queries for SLO checks.
 - Use Helm values to parameterize canary weights.
 - **Windows (drivers/services):**
 - Use site slots or blue/green server pools; ARR to route traffic for canary.
 - Validate driver signing & digital signature checks pre-deploy.
-

6) Security & Compliance

- SBOM created for every artifact and scanned for vulnerabilities; critical CVEs block or force Blue-Green + manual approval.
 - Access control: only pipeline tokens (with least privilege) can promote artifacts to production.
 - Audit: all decisions (`risk-decision.json` , `verdict.json` , playbooks`) stored with build metadata for compliance.
-

7) Observability & Dashboards (must-haves)

- **Release Compare:** candidate vs baseline (error_ratio, p95, cpu, mem, KPI)
 - **Canary Progress:** current weight, windows passed, verdict timeline
 - **Rollback Incidents:** list of automated rollbacks with reasons and links to artifacts
 - **ML/Policy Health:** shadow vs enforced policy performance (false positives, false negatives)
-

8) Operational Playbooks (short list)

- How to run a manual Blue/Green cutover (k8s helm swap / LB switch / Windows binding swap) and how to rollback.
 - How to run the Analyzer in shadow vs enforced mode and tune thresholds.
 - How to quarantine a firmware artifact and revert fleet if needed.
-

9) Storage & Retention

- Artifacts: retain candidates for N days (configurable); keep previous production artifacts for rollback.
 - Metrics: Prometheus retention policy (e.g., hot store 15d, long-term 90d in remote storage) to support baseline comparisons.
 - Logs & Traces: retain per compliance needs.
 - ML Data Lake: store parquet datasets for model training (retention per data policy).
-

10) Rollout Roadmap (phased)

1. **Phase 0:** Telemetry hygiene, per-OS environment parity, artifact metadata standardization.
 2. **Phase 1:** Implement CI per Step 2, publish candidate artifacts, create decision stage (heuristic, shadow).
 3. **Phase 2:** Integrate Orchestrator + Analyzer gates (EWMA/z-score), run shadow for 2–4 weeks.
 4. **Phase 3:** Enable enforcement for pilot services; implement auto-rollback playbooks.
 5. **Phase 4:** Build feature store + ML models; transition risk scoring to model-backed decisions (shadow → canary → enforced).
-

11) Owners & RACI (summary)

- Platform (owner): CI pipelines, orchestration, rollback controllers
 - Observability (owner): instrumentation, Prometheus/Loki/Tempo, dashboards
 - SRE (owner): runbooks, on-call, incident triage
 - Dev teams (owner): repo changes, tests, SBOMs, supporting telemetry
 - Security (owner): SBOM scanning & approval
 - Data/ML (owner): feature store and models
-

12) Immediate 30-Day Plan (practical next steps)

1. Standardize `build-info.json` and `risk-inputs.json` and ensure all repos emit it. (week 1)
 2. Deploy Prometheus/Loki/Tempo (or confirm access) and instrument 5 pilot services. (week 1–2)
 3. Add `Decide Strategy` stage in Jenkins in shadow mode, record `risk-decision.json` to Artifactory. (week 2)
 4. Build Release Compare dashboards and configure EWMA alerts. (week 2–3)
 5. Run shadow runs and refine thresholds, then enable enforcement for 1–2 pilot services. (week 3–4)
-

Closing

This architecture balances safety and automation: it gives a clear path from reproducible multi-OS builds to AI-driven deployment decisions, progressive delivery, fast detection and rollback, and a learning loop to continuously improve. If you want, I can now export a printable diagram (SVG/PDF) and a one-page operational runbook for the Release Decision step. Which would you like next?