

AI-Driven CI/CD — Four Methodologies

Scope: Procedure-only, step-by-step implementation plans derived from our previous design for a Jenkins + Multi-OS (ARM/Linux/Windows) pipeline with progressive delivery, anomaly detection, auto-rollback, and continuous learning. **No code** included.

Methodology 1: AI-Driven Adaptive Deployment Architecture (Jenkins + Multi-OS)

Objectives

- Automatically select **canary / blue-green / rolling** strategy per change.
- Orchestrate deployments across **ARM, Linux (K8s), Windows (IIS/services)**.
- Capture decision evidence and publish to **Artifactory Build-Info**.

Inputs

- Git change metadata (diff, LOC, files, deps, coverage delta)
- Historical success/rollback stats
- Current traffic/load & business calendar

Outputs

- Chosen strategy + rollout pace (% weights or batches)
- Deployment status (promoted/paused/rolled back)
- Evidence bundle (metrics snapshots, decisions, links)

Procedure (Implementation Steps)

1. **Establish Environments & Targets**
2. Define prod/stage clusters (K8s), Windows pools/sites, ARM cohorts.
3. Catalog services and OS support matrix.
4. **Define Strategy Policies**
5. Thresholds for **risk→strategy** mapping (Rolling < Canary < Blue/Green).
6. Allowed time windows; change freeze rules.
7. **Collect Decision Signals**
8. Standardize metrics labels (RED/USE) and user KPIs.
9. Normalize Jenkins build/test outcomes and dependency changes.
10. **Risk Scoring (Heuristic v1)**
11. Weight LOC, dep bumps, critical modules, past failure rate, coverage delta.
12. Produce `risk_score` 0–100; persist with the build record.
13. **Select Strategy & Pace**
14. Map score→strategy; define pace schedule (e.g., 10→25→50→100 or Blue/Green cutover gate).
15. **Wire Orchestrators**

16. K8s: prepare Helm/Argo Rollouts/Flagger specs.
17. Windows: prepare Blue/Green slots, ARR/LB weights.
18. ARM: define OTA cohorts and promotion rules.
19. **Gate with Observability**
20. Pre-deploy smoke checks; ensure metrics/alerts present for each service.
21. **Execute & Record**
22. Run the selected strategy; record decisions, windows, SLO status.
23. **Finalize & Publish**
24. On success: promote to 100%, tag **latest** in Artifactory; on failure: mark rollback and freeze.
25. **Review & Tune**
26. Weekly review of risk thresholds vs. outcomes; adjust mappings and paces.

Success Criteria

- <2% unexpected rollbacks; >30% reduction in mean deploy time; full decision evidence attached to artifacts.
-

Methodology 2: Real-Time Monitoring with Anomaly Detection

Objectives

- Detect regressions during rollout in **minutes**.
- Produce a **verdict**: promote / slow / rollback.

Inputs

- Metrics (error rate, p95/p99 latency, CPU/memory, throughput)
- Logs (error signatures, new patterns)
- Traces (error spans, slow endpoints)
- Baseline from last stable release (same time-of-day if seasonal)

Outputs

- Window-by-window health status
- Anomaly score + SLO compliance
- Action recommendation (promote/slow/rollback) with reasons

Procedure (Implementation Steps)

1. **Standardize Telemetry**
2. Instrument apps with OpenTelemetry; ensure consistent labels.
3. Enable exporters for K8s (kube-state, cAdvisor), Windows (IIS/Win exporters), ARM beacons.
4. **Define SLOs & Windows**
5. Set service SLOs (e.g., error<1%, p95<300ms); choose window length (1–2m) and required consecutive passes.
6. **Create Baselines**
7. Snapshot metrics for the last stable version; align to comparable load periods.

8. Recording Rules & Alerts

9. Create summarized series (error ratio, latency quantiles, CPU/mem utilization).

10. Anomaly Methods (Phase-in)

11. Phase 1: thresholds + EWMA smoothing.

12. Phase 2: baseline z-scores; seasonal decomposition (STL).

13. Phase 3: multi-metric outlier detection (Isolation Forest) for hard drifts.

14. Decision Policy

15. Encode promote/slow/rollback based on SLO + anomaly results and risk level.

16. Integrate with Rollouts

17. Connect analysis steps to rollout gates (K8s AnalysisTemplates, Jenkins stages, Windows checks).

18. Evidence & Notifications

19. Store verdicts, plots, and links; notify Slack/Jira with context and runbook.

20. Shadow Period

21. Run detection in observe-only mode; measure false positives/negatives.

22. Enforce

23. Activate gating once shadow metrics are acceptable; review monthly.

Success Criteria

- Median detection time < 5m; false-positive rate < 10%; no undetected critical regressions.
-

Methodology 3: Auto-Rollbacks & Self-Healing

Objectives

- Reduce blast radius via **fast, safe reversion**.
- Automatically fix common infra/app issues without human intervention.

Inputs

- Anomaly verdicts; SLO breaches; health checks
- Catalog of **idempotent playbooks** per platform (K8s, Windows, ARM)
- Last-known-good versions and configs

Outputs

- Completed rollback actions (version/traffic/capacity)
- Executed self-healing actions with status
- Audit trail and freeze conditions when necessary

Procedure (Implementation Steps)

1. Define Rollback Policies

2. Criteria by stage (% traffic) and severity; freeze rules post-rollback.

3. Catalog Playbooks

4. K8s (abort/promo to stable, helm rollback, restart, HPA/VPA), Windows (slot swap, ARR weight, AppPool restart), ARM (cohort revert, exclude devices).

5. Establish Safety

6. Idempotency, locks per service, exponential backoff, and caps on retries.

7. Wire Controllers

8. Create discrete rollback and self-healing controllers invoked by CD gates.

9. Config & Schema Safety

10. Maintain LKG configs; apply expand/contract for DB migrations; use feature flags for risky paths.

11. Verification

12. Post-action checks (health endpoints, KPIs) before unfreezing or re-promoting.

13. Evidence & Comms

14. Persist decisions, timestamps, affected cohorts/nodes; auto-open incident tickets with links.

15. Chaos & Drills

16. Regular game days to validate automation; update runbooks from findings.

17. Governance

18. Review rollbacks weekly; prune ineffective playbooks; refine thresholds.

Success Criteria

- MTTR reduction > 50%; rollback correctness ~100%; no repeated incident without runbook update.
-

Methodology 4: Continuous Learning from Logs → Pipeline Efficiency

Objectives

- Turn operational data into **policy improvements** that speed up builds, tests, and safer deployments.

Inputs

- Jenkins build/test outcomes, durations, cache hits
- Deploy decisions & outcomes (promote/slow/rollback)
- Observability metrics; log templates; trace errors
- Artifact metadata & security scans

Outputs

- Updated policies: risk thresholds, test/build selection, rollout pacing
- Flaky test quarantine lists; cache optimization guidance
- Governance reports and dashboards

Procedure (Implementation Steps)

1. Ingestion & Normalization

2. Consolidate logs/metrics/build data into a unified schema (parquet in object store).

3. Feature Store Setup

4. Create offline/online features (LOC, dep bumps, fanout, flakiness, cacheability, early anomaly scores).

5. Define Targets & KPIs

- 6. Predictors for failure/rollback; objectives for time saved, precision/recall, false-skip caps.
- 7. **Train & Evaluate Policies**
- 8. Train models (risk, test selection, build skip, pacing); evaluate vs rolling baselines.
- 9. **Shadow Policies**
- 10. Run counterfactual simulations; compare to status quo without affecting prod.
- 11. **Controlled Rollout**
- 12. Canary the policies on subset of repos/services; enable rollback to prior policy version.
- 13. **Governance & Versioning**
- 14. Track experiments and register policies; store lineage linking builds, data, and decisions.
- 15. **Feedback Integration**
- 16. Feed improved policies back into Methodologies 1–3; schedule periodic retraining.
- 17. **Reporting**
- 18. Dashboards for CI time saved, defect detection efficiency, rollback rate, canary dwell time, cost per change.

Success Criteria

- $\geq 25\%$ CI time saved; $\geq 50\%$ less flaky test noise; stable or reduced rollback rate; positive cost trend.
-

Cross-Cutting Readiness Checklist

- ☐ Metrics coverage for every service & environment
 - ☐ Baseline references for last stable release per service
 - ☐ Clear SLOs and alert routes with runbooks
 - ☐ Risk→Strategy policy file versioned in Git
 - ☐ Idempotent rollback & healing playbooks tested in staging
 - ☐ Evidence pipeline to Artifactory/DB + notification hooks
 - ☐ Weekly review cadence and owners for thresholds/models/policies
-

RACI Snapshot (Who does what)

- **Owners:** Platform/DevOps team (policies, controllers, observability)
 - **Contributors:** Service teams (SLOs, KPIs, coverage mapping)
 - **Reviewers:** SRE & Security (gates, runbooks, failure classes)
 - **Approvers:** Engineering leadership (risk thresholds, freeze policies)
-

Implementation Roadmap (Phased)

- **Phase 0 (2–3 weeks):** Telemetry hygiene, SLOs, environment matrix, artifact tagging.
- **Phase 1 (3–4 weeks):** Adaptive strategy (heuristic), anomaly thresholds, evidence pipeline.
- **Phase 2 (4–6 weeks):** Auto-rollback playbooks, shadow anomaly, controlled enforcement.
- **Phase 3 (6–10 weeks):** Continuous learning loop (shadow → canary → full), flaky/test selection, build skip predictor.

- **Ongoing:** Monthly policy review, chaos drills, governance reports.